

# Manual of EImageFit

Last update: December 2012

Author: Javier Velázquez-Muriel

Last fully tested revision: SVN 15319

## 1) Introduction

EImageFit is an application to build models of macromolecular assemblies using restraints from EM images (class-averages). It is based on the paper:

*Velázquez-Muriel J, et al. (2012) Assembly of macromolecular complexes by satisfaction of spatial restraints from electron microscopy images. Proc Natl Acad Sci U S A 109(46):18821-18826.*

Apart from EM images, the method relies on other types of restraints:

- Maximum distance restraints imposing a maximum distance between a pair of residues, as derived from cross-linking experiments. You can of course impose maximum distance between any part of the complex. As shown in the paper, this type of restraints are required to accurately determine the orientation between the subunits of an assembly.
- Proximity restraints enforcing a pair of subunits to be within a certain distance, suitable for proteomics data.
- Excluded volume restraints preventing the components of the assembly from interpenetrating.
- Geometric complementarity restraints favoring large contact surfaces between interacting subunits. These restraints can be somewhat useful if the subunits of the assembly are expected to have large contact surfaces. Among all of the restraints considered here, they are the least useful.

EImageFit samples for good conformations of the macromolecular assembly using multiple molecular docking (optional, but it can be very useful), Simulated Annealing Monte Carlo optimization, and sampling with the discrete optimizer DOMINO. It is also straightforward to incorporate other restraints to the method.

## 2) Installation and requirements

EImageFit is part of IMP, and therefore has all the dependencies that IMP has. But you are aware of it if you are reading this. On top of them, the *em2d* module, which is the base of this application, has some other requirements. The list is intimidating, but you'll see that it makes

sense:

- **Python2.7.** It is probably on your machine already. It is required for the management of results, as they are stored as a SQLite database.
- **openCV.** openCV is a library for computer vision that provides general functionality for dealing with images. EImageFit uses it for speed and support of graphical formats for the images. We have tested that EImageFit works with OpenCV 2.1, 2.2 and 2.3. Installing openCV should not be very traumatic. Here is some quick instructions that work for us:

On linux:

- Download the source code from <http://opencv.itseez.com/> and put it in a directory, `opencv_install` for example. Go to this directory.
- Call `cmake` to build the library. There is only one thing to pay attention to, the directory where you want to install the program:  
`cmake -D CMAKE_BUILD_TYPE=RELEASE -D CMAKE_INSTALL_PREFIX=/home/you/OpenCV-2.3.0/ -D BUILD_PYTHON_SUPPORT=ON ..`
- In our case, one more trick was necessary. We had to go to the file: `MakeCache.txt` and change `PYTHON_EXECUTABLE:FILEPATH=/usr/bin/python2.3` for `PYTHON_EXECUTABLE:FILEPATH=/usr/bin/python2.7`
- Execute `make`
- Execute `make install`
- Now you should have openCV in the directory `/home/you/OpenCV-2.3.0/`. You may have some other problems, and here is the point where you will want to go to the openCV documentation and fight a bit with it.

On Mac:

MacPorts, a tool to get Unix programs working on Mac, provides a version of openCV. Installing openCV should be really easy. Learn a bit about the "port" command from macports, and install openCV. It should be only a line.

```
sudo port install opencv
```

Remember that as any other dynamic library, you must have its path in the environment variable `LD_LIBRARY_PATH`. In our example, the path to include is `/home/you/OpenCV-2.3.0/lib`. In *bash* shell, the command to set the `LD_LIBRARY_PATH` variable is:  
`export LD_LIBRARY_PATH=/home/you/OpenCV-2.3.0/lib`

- **The GNU scientific library** <http://www.gnu.org/software/gsl/>. EImageFit uses it because it has a good implementation of the Simplex method for function optimization. If you have an standard linux distribution, most probably it is already installed. The gsl module in IMP should then take care of the rest. For mac, as simple as before:  
`sudo port install gsl`
- **The docking program HEXDOCK.** <http://hex.loria.fr/>. The program is optional but very good to have. EImageFit uses it to do docking between subunits that are related by cross-links. But it is worth mentioning that EImageFit can also work with any other docking program, or simply using the cross-linking restraints and no docking at all. HEXDOCK is not required for building neither the IMP em2d module nor EImageFit.

If you have all these dependencies, IMP will be ready to compile the em2d module. The em2d module depends on other IMP modules: core, atom, em, gsl, container, and em. Make sure that you compile them by setting the proper variables in your IMP configuration file. All the required modules have their own dependencies, but at the time this manual was written you should be fine with the libraries mentioned above.

### 3) Input data

You need only three things to get models using EMageFit:

1. A set of PDB files with the components of your assembly.
2. A set of EM images.
3. A configuration file.

Let's describe them a little bit.

**The PDB files with the components of the assembly.** Each PDB file must contain a protein, a DNA strand, or a subcomplex. It is possible to have various different chains within the same PDB, and all of them will be considered as a rigid body. All the chains that are going to be assembled must have different ID. i.e., you cannot have a chain ID A in two different files. All the atoms in the PDB file will be used. If you don't want to have duplicated atoms, due to PDB records as ANISOU, you need to remove them. We tend to be more radical and remove other records like REMARK, SOURCE, COMPND, SEQRES, DBREF, CONECT. They are not relevant for the type of problem that EMageFit solves.

**The EM images.** IMP can understand 3 image formats. Spider, JPG and TIFF. The format that you want to use 99% of the time is Spider, as it is specific for EM. Of course there are other formats for EM, and you can use the free program [em2em](#) to convert them to Spider. Conversion to JPG and TIFF are nice to have as more traditional graphic formats. There is a script in the directory of this manual to do conversions.

Each EM image has to be a separate file. The images that you want to use need to be listed in a "selection file". This is just a file with 2 columns, with the name of the images and 0/1. If there is a 0, that image is not used. For example, a selection file [myselection.sel](#) like this:

```
image1.spi 1
image2.spi 0
image3.spi 1
```

means that you have 3 images but only want to use image1.spi and image3.spi for modeling.

**A configuration file.** A configuration file is just a Python file with classes that describe all the parameters and restraints that you want to use. Using a python file as configuration file makes adding new parameters to your simulation trivial. It is better to describe the file in situ, so open the file [config\\_example.py](#) located in the example\_3sfd directory to find a detailed description of each of the parameters. Once that you understand all options, you can of course remove all the lengthy comments.

## 4) How to get models

Getting models requires 4 steps:

1. Doing the preliminary dockings.
2. Obtaining models with Monte Carlo optimizations.
3. Gathering the solutions from the Monte Carlo optimizations.
4. Combining models from Monte Carlo with DOMINO to get even better models.

### 4.1 Pairwise dockings

EMageFit performs docking between components that are subject to cross-linking restraints using the program HEXDOCK. This is what EMageFit does:

- Finds a very rough estimation of the orientation between two components by minimizing the distance between the aminoacids implied in the cross-linking restraints. Admittedly, this is not going to be very good, but it will help the HEXDOCK program providing a "hint" of the orientation. You then will tell HEXDOCK that you don't want to search all possible orientations for the ligand, just a given angle around the orientation obtained from the rough guess. In our experience, it works well in many cases.
- Determines an optimal way of doing the required dockings. It finds the component of the assembly that needs to be kept anchored, and establish an order for the dockings. The docking order is based on the maximum spanning tree of the graph built by the component connections. Here is an example: Let's say that you have a complex with 5 chains: A B C D E, and your cross-linking restraints say that the components should be connected like this:

A-B-D-E

  \|/  
  C

The weight for an edge is the number of restraints between the components that it connects. After computing the maximum spanning tree, you could a graph similar to this:

A-B-D-E

  |  
  C

Which says that B should be anchored and be the first receptor, as it is the component with the largest number of neighbors. And edge indicates that a docking should be done. In our case:

  A (ligand) docked to B (receptor)  
  C docked to B  
  D docked to B  
  E docked to D

- Runs HEXDOCK. Optionally, you can get the docking order computed in the previous step and use your favorite docking program/server to do these dockings. You will need to recover from your docking program the relative transformation of the ligand respect to the receptor for each solution. See the explanation for the [emagefit\\_dock.py](#) script to see how you could integrate your favorite program.

- Filters the docking solutions that are compatible with the cross-linking restraints. As an emergency measure, if there are no solutions compatible with the restraints, all of them are taken. Of course this implies the risk of using solutions that are not very accurate. The rough estimation calculated during the first step is also kept.

The entire procedure described above can be done with the command:

```
imppy.sh python emagefit.py --exp config_step_1.py --dock --log file.log
```

Now you have to take the information from the dockings and put it in the configuration file. You need to indicate which component is anchored and provide the files of relative transformations from the dockings. The options to modify are `self.anchor` and `self.dock_transforms`. The example of section 6 will show you how.

As mentioned before, the pairwise dockings are optional. EMageFit can work without them. In that case all the Monte Carlo moves will be random. You should indicate that you don't have docking solutions by setting the option `self.non_relative_move_prob` (probability of doing a not docking-related move) to value 1. This means that EmageFit will always do a random move. Another strategy is possible: If the docking program does not produce solutions compatible with the cross-linking restraints you may still want to use them. They have reasonable conformations after all, without classes, and perhaps not far from some other conformation that actually satisfies the restraints. It is advisable then to set `self.non_relative_move_prob` to a low number (say 0.2, meaning that a move from a docking solution will be chosen only 20% of the time).

The option `--log` is not mandatory. If you use it (recommended), you will get a file with the information for the modeling. Otherwise, all the information will be printed on the screen. This logging information is coming only from the python interface of this application, and is different from the IMP logging system. You can select the granularity of the logging by using the variables from the Python *logging* module: DEBUG, INFO, etc.

## 4.2 Obtain models with Simulated annealing Monte Carlo optimization

Once that the relative docking transformations are set, you have to do Monte Carlo optimizations for getting models. The command is:

```
imppy.sh python emagefit.py --exp config_step_2.py --o montecarlo1.db --log file.log
--monte_carlo -1
```

The `--monte_carlo` option informs `emagefit.py` how to handle the random seed at the beginning of the Monte Carlo optimization. When using the script in a cluster of computers and sending multiple jobs at the same time, the random seed obtained when using the machine time is the same for all jobs. To avoid that situation, you can pass an integer. If the number is -1, the machine time is used as the seed (useful for one core).

What you get after the optimization is a SQLite database (`montecarlo1.db` in this example) with only one solution. This is so because the idea is to run `emagefit.py` with Monte Carlo as many times as models you want. In a computer cluster all this can be done in parallel. The result is a

set of database files: `montecarlo1.db`, `montecarlo2.db`, etc.

### 4.3 Gather the results of all Monte Carlo optimizations

Now is time to put all the Monte Carlo solutions together. It is done with the command:

```
imppy.sh python emagefit.py --o all_montecarlo.db --gather {all database files}
```

Where {all database files} means the name of all the files to join. Something like `montecarlo*.db`, if you decided to use `montecarlo1.pdb`, `montecarlo2.pdb`, etc. as the names of the databases.

### 4.4 Combine the models from Monte Carlo with DOMINO

The solutions in `all_montecarlo.pdb` are already solutions for the modeling. They are a set of discrete solutions that can be improved by combining the positions of the components in all of them. For example, if you have 100 solutions from the Monte Carlo experiments, then you have 100 possible positions for each component. The positions should be already correct, but what you can achieve with DOMINO is to explore all the possible combinations Monte Carlo solutions further improving their quality. If the assembly that you want to model has 4 components, using DOMINO you are exploring the  $100^4$  possible combinations. To do the task, first you have to go the configuration file and change the value of the member variable `self.read` in the `DominoSamplingPositions` class. The value that you want is the name of the database of Monte Carlo results. For example, `all_montecarlo.db`. Then run the command:

```
imppy.sh python emagefit.py --exp config_step_3.py --o domino_models.db --log file.log
```

This will produce a database `domino_models.db` with all the results.

## 5) Visualizing the models and understanding the information in the database of solutions

The database of results contains all the positions for the rigid bodies in the solutions. To write some of these solutions, the command is:

```
imppy.sh python emagefit.py --exp config_step_3.py --o domino_models.db --w 10 --orderby em2d --log file.log
```

In this case the option `--o` does not modify the database, only uses it for reading the positions of the components. The option `--w` says that we want 10 models. The option `--orderby` is the name of the restraint used to sort the models. When using `--orderby em2d` in the example above, you say that you want the 10 best models according to the value of the `em2d` restraint. Another typical value is `total_score`. The solutions are written to the files `solution-000.pdb`, `solution-001.pdb`, and so on. Each record for a solution in the database contains the following

information:

1. **Solution\_id** - A unique number that identifies the model. Note: solution\_id=0 does not mean the best solution. It is only an identifier.
2. **assignment** - Is the set of numbers identifying a combination in domino. For the previous example with 4 components an 100 positions, one assignment could be "11|23|45|76", and the meaning is: "Use position of the first component in solution 11, combined with position the second component in solution 23, etc".
3. **Reference frames**. These are the values used to build an [algebra.ReferenceFrame3D](#) object in IMP. There is one reference frame per component of the assembly. Generating a solution is as simple as setting the reference frame of each of the rigid bodies of the components of the assembly.
4. **Total\_score** - The total value of the scoring function.
5. **{restraints}** this is a list of values for the restraints. There is one column in the database for each restraint. The list changes with the number and nature of the restraints. You can print their names by using the script [Database.py](#) stored in the directory [pyext/src](#) of the em2d module.

Using a database output file is very powerful, as you can query the data and decide how you want the information. On the other side, it is less comfortable than a simple text file. To help with that, see the file [quick\\_sql\\_query.py](#) in the directory of the example, which contains some typical SQL queries.

## 6) A complete example

Here is an entire example for one of the experiments in the paper. It is the modeling of the structure with PDB ID 3sfd. You can find it in the subdirectory [example\\_3sfd](#). The inputs are the files 3sfdA.pdb, 3sfdB.pdb, 3sfdC.pdb, 3sfdD.pdb and the images in the directory [em\\_images](#). The selection file is [images.sel](#). To do the dockings, run

```
imppy.sh python emagefit.py --exp config_step_1.py --dock --log dock.log
```

Open the file [dock.log](#) and search for the line:

```
INFO:buildxlinks:The suggested order for the docking pairs is [('3sfdB', '3sfdA'), ('3sfdB', '3sfdC'), ('3sfdB', '3sfdD'), ('3sfdD', '3sfdC')]
```

This line is telling you the dockings required, and the order recommended. The pairs are (Receptor,Ligand), so you need to dock 3sfdA into the 3sfdB, 3sfdC into 3sfdB, 3sfdD into 3sfdB, and 3sfdC into 3sfdD. If you have HEXDOCK and everything went well, you'll see a lot of new files. We have included them in the [outputs](#) directory too.

Files with the ligand in the position of the rough estimation of the orientations based on the cross-links:

```
3sfdB-3sfdD_initial_docking.pdb
3sfdB-3sfdA_initial_docking.pdb
3sfdD-3sfdC_initial_docking.pdb
3sfdB-3sfdC_initial_docking.pdb
```

Files with the first solution found by HEXDOCK:

```
3sfdB-3sfdA_hexdock.pdb
3sfdB-3sfdC_hexdock.pdb
3sfdB-3sfdD_hexdock.pdb
3sfdD-3sfdC_hexdock.pdb
```

Files of all the transformations of the ligand from HEXDOCK:

```
hex_solutions_3sfdB-3sfdA.txt
hex_solutions_3sfdB-3sfdD.txt
hex_solutions_3sfdB-3sfdC.txt
hex_solutions_3sfdD-3sfdC.txt
```

Files with the filtered solutions:

```
hex_solutions_3sfdB-3sfdA_filtered.txt
hex_solutions_3sfdB-3sfdC_filtered.txt
hex_solutions_3sfdB-3sfdD_filtered.txt
hex_solutions_3sfdD-3sfdC_filtered.txt
```

Files with the relative transformations of the ligand respect to the receptor:

```
relative_positions_3sfdB-3sfdA.txt
relative_positions_3sfdB-3sfdD.txt
relative_positions_3sfdB-3sfdC.txt
relative_positions_3sfdD-3sfdC.txt
```

The files with the relative transformations are the files that we want. The component to anchor is 3sfdB, because in the component with most neighbors. It is easy to identify because it is the first receptor in the list of docking pairs. If you don't have HEXDOCK You can specify that in the configuration file with `self.have_hexdock = False`. You will still get the order suggested, but you have to do the dockings with your favorite program. You have to compute the relative orientations of the ligand respect to the receptor and translate them into IMP transformations. If you can do that, then writing a file like `relative_positions_3sfdB-3sfdA.txt` is not difficult. Assuming that you have found the relative transformation given by 3 Euler angles ZYZ (phi, theta, psi) and a translation (x,y,z), here is the set of commands in IMP that will give you the transformation:

```
R = IMP.algebra.get_rotation_from_fixed_zyz(phi, theta, psi)
q = R.get_quaternion()
```

What you see in each line of `relative_positions_3sfdB-3sfdA.txt` is just:

```
q[0] | q[1] | q[2] | q[3] | x | y | z
```

You can also use the cross-linking restraints without relying on any docking solution. To do that, delete the `self.dock_transforms` option from the configuration file. Once that you have the dockings and the anchored component, go to the configuration file `config_step_1.py` and fill the values for `self.anchor` and `self.dock_transforms`:

```
self.anchor = [False, True, False, False]
self.dock_transforms = [
    ["3sfdB", "3sfdA", "relative_positions_3sfdB-3sfdA.txt"],
    ["3sfdB", "3sfdC", "relative_positions_3sfdB-3sfdC.txt"],
    ["3sfdB", "3sfdD", "relative_positions_3sfdB-3sfdD.txt"],
    ["3sfdD", "3sfdC", "relative_positions_3sfdD-3sfdC.txt"],
]
```



The file `config_step_2.py` contains all the changes.

The next step is running a Monte Carlo optimization. You can adjust the profile of temperatures, number of iterations, cycles, maximum displacement and angle tolerated for the random moves, and also the parameter `self.non_relative_move_prob`. This parameter indicates the probability for a component of doing a random move instead of a relative move. If you put 0.4 it means that the component acting as ligand (e.g. 3sfdA) will prefer to do a random move respect to its receptor (3sfdB) 40% of the time. A value of 1 ignores all the relative positions, and a random move is always chosen. The same applies to all other pairs of components. To run the optimization:

```
imppy.sh python emagefit.py --exp config_step_2.py --monte_carlo -1 --log
monte_carlo.log --o mc_solution1.db
```

Probably you noticed that there were changes in the MonteCarloParams of `config_step_2.py` respect to `config_step_1.py`. The new parameters for Monte Carlo were set to get a very short simulation. You'll get a garbage model (quickly). The actual parameters used during the benchmark for the paper are those on `config_step_1.py`. Once the script has finished, there should be two new files in the directory, the logging file and the database with the result:

```
monte_carlo.log
mc_solution1.db
```

After obtaining a set of Monte Carlo models, the command to gather all the individual solution files into a single one is:

```
imppy.sh python emagefit.py --o monte_carlo_solutions.db --gather mc_solution*.pdb
```

For this example we have included the file `monte_carlo_solutions.db`, which contains the results of 500 Monte Carlo runs.

The last part of the modeling is running DOMINO employing the configuration file `config_step_3.py`. This new configuration file has changes respect to `config_step_2.py` in the classes `DominoSamplingPositions` and `DominoParams`. The parameters are:

- **self.read** is the file with the Monte Carlo solutions obtained before.
- **self.max\_number** is the maximum number of solutions to combine. In this example, with 500 solutions and 4 components, we would have to explore  $500^4$  combinations, which is unfeasible. This number allows you to reduce that. In the example, we set 5, and therefore only  $5^4$  combinations are explored.
- **self.orderby** is the name of the restraint used to sort the Monte Carlo solutions. Here the value is "em2d", so the best 5 solutions according to the em2d score will be combined with DOMINO. You could try using "total\_score" too.
- **self.heap\_solutions**. This is a rather technical parameter. It is the number of solutions that you keep each merging step in DOMINO. The larger the number, the better the space of  $5^4$  combinations is explored, at the cost of a larger running time. Here we put 200.

The command is:

```
imppy.sh python emagefit.py --exp config_step_3.py --log domino.log --o domino.db
```

You will get files `domino.log` and `domino.db`, containing the logging and the database of solutions, respectively. Once again, the parameters used for domino were selected to get a quick answer. For the example, we have included the file `domino_solutions.db`, which we obtained during our benchmark. We used `self.max_number=50` and `self.heap_solutions=2000`.

To write some solutions run:

```
imppy.sh python emagefit.py --exp config_step_3.py --w 10 --o domino.db --orderby em2d
```

The solutions will be in the files `solution-*.pdb`. They will not be very good, because they came from sampling only  $5^4$  combinations. But try the file `domino_solutions.db` that we obtained during our benchmark:

```
imppy.sh python emagefit.py --exp config_step_3.py --w 10 --o domino_solutions.db --orderby m2d
```

## 7) Individual description of the scripts

**emagefit.py.** This script can be used for all the stages of modeling: Docking, Monte Carlo optimization, gathering of solutions from the Monte Carlo runs, DOMINO sampling, and finally write the resulting models.

**emagefit\_dock.py** - A wrapper for the program HEXDOCK. It uses HEXDOCK in text mode to perform a docking of a subunit (the ligand) into another subunit (the receptor). The script can be used as a standalone program to perform a docking or write the solutions. The script can be modified to use any docking program, by doing a couple of changes:

- The class `HexDocking` is called only from `emagefit.py` and uses the `dock()` method. You can change `HexDocking` for another wrapper class providing a `dock()` method that saves the results to a file `fn_transforms`.  
`emagefit.py` also calls the functions `read_hex_transforms()` and `filter_docking_results()`. You only need to adapt the function `parse_hex_transform()`, which both of them use, to your docking program.

**emagefit\_cluster.py.** Performs clustering of the solutions stored in a database file. You can run it as standalone program. The help of the script gives the parameters required, and a typical command is:

```
imppy.sh python emagefit_cluster.py --exp config_step_3.py --db domino_solutions.db -o clusters.db --n 100 --orderby em2d --log clusters.log --rmsd 10
```

To write the elements of the first cluster:

```
imppy.sh python emagefit.py --exp config_step_3.py --o domino_solutions.db --wcl clusters.db 1
```

**emagefit\_score.py.** This script returns the em2d score for a model using the EM images. It is useful for comparing models obtained by other sampling algorithms apart from the one described in the paper. This script supersedes the scripts `em2d_score` and `em2d_single_score` that you will find in the `build/bin` directory. To score a model the parameters required are:

- The PDB file of the model.
- The selection file for the EM images.
- The pixel size of the EM images.
- The number of projections used for the coarse registration step of the scoring.
- The resolution used to generate the projections. The model is downsampled to this value of the resolution before projecting it. The larger the value, the blurrier the projections generated. For our benchmark we used a value as low as 2, as the results are not very different from using lower resolutions. For images of poor quality, with no distinguishable features, values of 10-15 may be used.
- Images per batch. This parameter is used to avoid running out of memory when the number of images used for scoring a model is large. The scoring is done keeping in memory only the number of images specified by the parameter.
- An example:

```
imppy.sh python emagefit_score.py structure.pdb myimages.sel 3.6 20 5 100
```

**convert\_spider\_to\_jpg.py.** Does what it says.

Other scripts can be found in the directory [pyext/src](#) of the em2d module:

**buildxlinks.py** - Contains all the code for generating the order of the dockings. It also contains the class `InitialDockingFromXlinks`, which is used to move the position of the subunits acting as ligand close to the receptor.

**DominoModel.py.** Contains the `DominoModel` class, which has a `IMP.Model` as the main member. The class manages the details of setting the model restraints, performing the Monte Carlo runs, configuring the DOMINO sampler, and storing the results in a database.

**MonteCarloRelativeMoves.py.** Contains the class `MonteCarloRelativeMoves` for setting and configuring a simulated annealing Monte Carlo optimizer. It also manages the profiles of temperature and iterations for the sampling. The optimizer uses one `em2d.RelativePositionMover` object per docking to propose relative moves of a ligand respect to the receptor.

**restraints.py.** Creates the restraints used for the modeling. It is called from `DominoModel.py`.

**sampling.py.** The script allows you to set the positions and orientations for the components of the assembly before combining them using DOMINO. In the paper we ended using the set of Monte Carlo solutions, but you can use the script to set any other combination of positions and orientations for the the subunits.

**solutions\_io.py.** Contains the class `ResultsDB` for managing the database of solutions obtained during modeling.

**Database.py**, **argminmax.py**, **csv\_related.py**, and **utility.py** are supporting scripts. **ResultsDB** inherits all the basic functionality from **Database.py**, a wrapper for SQLite databases. The wrapper is easy to use, general, and it does not depend on IMP, so it may be useful for managing your data too.

Some other scripts are stored in the directory [pyext/src/imp\\_general](#). These scripts are general and perform basic and/or frequent tasks in IMP. They can be helpful for your own IMP scripts:

**representation.py**. The main script. It contains functions for obtaining the representation of an assembly from one or more PDB files, creating rigid bodies for the components of the assembly, simplifying the structure of a protein using beads, getting coordinates and distance between residues, etc.

**alignments.py**. A couple of functions to align assemblies.

**comparisons.py**. Functions to compute the cross-correlation coefficient between density maps, RMSD and DRMS between models, and placement score for the subunits of an assembly as defined in the paper.

**movement.py**. Functions for transforming a rigid body or a structure.

If you have more questions, send us an email to [imp-dev@salilab.org](mailto:imp-dev@salilab.org) (for technical questions regarding the code or functionality), or [imp-users@salilab.org](mailto:imp-users@salilab.org) (for questions regarding the use of the script).

Happy integrative modeling!